

Chapter 3: Introduction to SQL

Practice Exercises:

- (3.1) Write the following queries in SQL, using the University Schema.
- Find the titles courses in the Comp. Sci. department that have 3 credits.
 - Find the ID's of all students who were taught by an instructor named Einstein, make sure there are no duplicates in the result.
 - Find the highest salary of any instructor.
 - Find all instructors earning the highest salary.
 - Find the enrollment, across all sections, in Autumn, 2009.
 - Find the maximum enrollment, across all sections in Autumn 2009.
 - Find the sections that had the maximum enrollment in Autumn 2009.

Soln:

- select title from course where dept_name = 'Comp. Sci.' and credits = 3
- select distinct student.ID from (student join takes using (ID)) join (instructor join teaches using (ID)) using (course-id, sec-id, semester, year) where instructor.name = 'Einstein'
- select max(salary) from instructor

(d) select ID, name from instructor
where salary = (select max(salary) from instructor)

(e) select course_id, sec_id, count(ID) from
section natural join takes where semester = 'Autumn'
and year = 2009 group by course_id, sec_id.

Another way:

(~~f~~) select course_id, sec_id, (select count(ID))

from takes where takes.year = section.year
and takes.semester = section.semester
and takes.course_id = section.course_id
and takes.section_id = section.section_id
from section where semester = 'Autumn' and
year = 2009 .

(~~g~~) select max(enrollment) from (select count(ID)
as enrollment from natural join takes
where semester = 'Autumn' and year = 2009
group by course_id, sec_id).

(g) with sec_enrollment as (

select course_id, sec_id, count(ID) as enrollment
from section natural join takes
where semester = 'Autumn' and year = 2009
group by course_id, sec_id)

select course_id, course_id from sec_enrollment
where enrollment = (select max(enrollment))
from sec_enrollment.

(3.2) Suppose you are given a relation grade-points (grade, points), which provides a conversion from letter grades in the takes relation to numeric scores; for example an "A" grade could be specified to correspond to 4 points, an "A-" to 3.7 points, a "B+" to 3.3 points, a "B" to 3 points, and so on. The grade points earned by a student for a course offering (section) is defined as the number of credits for the course multiplied by the numeric points for the grade that the student received.

Given the above relation, and our university schema, write each of the following query in SQL. You can assume for simplicity that no takes tuple has the null value for grade.

- Find the grade-points earned by the student with ID 12345, across all courses taken by the student.
- Find the grade-points average (GPA) for the above student, that is, the total grade points divided by the total credits for the associated courses.
- Find the ID and the grade point average of every student.

Soln :

(a) (select sum (credits * points))

from (takes natural join courses) natural join
grade-points (where ID = '12345')

union

(select * from student where takes.ID = '12345' and not exists (select * from takes where takes.ID = '12345'))

(b) (select sum(credits*x.points)/sum(credits) as
GPA from (takes natural join course) natural
join grade-points where ID = '12345')

Union

(select null as GPA from student where
takes.ID = '12345' and not exists (select * from
takes where takes.ID = '12345'))

(c) (select ID, sum(credits*x.points)/sum(credits)
as GPA from (takes natural join course) natural
join grade-points group by ID)

Union

(select ID, null as GPA from student where
not exists (select * from takes where takes.ID =
student.ID))

(3.3) Write the following inserts, deletes or
updates in SQL, using -the University schema.

(a) Increase the salary of each instructor in
the Comp. Sci. department by 10%.

(b) Delete all courses, that have never been
offered

(c) Insert every student whose tot_cred attribute
is greater than 100 as an instructor in the same
department with a salary of \$10,000.

(a) update instructor set salary = salary * 1.10
where dept-name = 'comp. sci'

(b) delete from course where course_id not
in (select course_id from section)

(c) insert into instructor select ID, name,
dept-name, 4000 from student
where tot-cred > 100.

(b.4) consider the insurance database of Figure
B.18 where the primary keys are underlined.
construct the following SQL queries for this
relational database.

- Find the total number of people who owned
cars that were involved in accidents in 1980
- Add a new accident to the database; assume
any values for required attributes.
- Delete the Mazda belonging to "John Smith".

Soln:

(a) select count(distinct name)
from accident, participated, person
where accident.report-number = participated.report-number
and participated.driver_id = person.driver_id
and date between date '1980-00-00' and
date '1980-12-31'

(b) Insert into accident values (4007, '2001-01-01', 'Berkeley')

insert into participated

select o.driver_id, c.license, 4007, 3000

from person p, owns o, car c where

p.name = 'Jones' and p.driver_id = o.driver_id and
o.license = c.license and c.model = 'Toyota'

(c) delete car where model = 'Nazda' and
license in (select license from person p,
owns o, where p.name = 'John Smith' and
p.driver_id = o.driver_id)

(3.5) Suppose that we have a relation marks
and we wish to assign grades to student
based on the score as follows: grade F if
score < 40, C if score < 60, grade B if
 $60 \leq \text{score} < 80$ and grade A if $80 \leq \text{score}$.

(a) Display the grade for each student, based on
the marks relation.

(b) Find the number of students with each grade.

Soln:

(a) select ID, case

when score < 40 then 'F'

when score < 60 then 'C'

when score < 80 then 'B'

else 'A'

end

from marks,

(b) with grade as (

select ID, case

when score < 40 then 'F'

when score < 60 then 'C',

when score < 80 then 'B'

else 'A'

end as grade)

from marks)

select grade, count(ID) from grades
group by grade.

(3.6) The SQL like operator is case sensitive, but the LOWER() function on strings can be used to perform case insensitive matching. To show how, write a query that finds departments whose names have the string "sci" as a substring. Regardless of the case.

SOLN: Select dept-name from department
where LOWER(dept-name) like '%sci%'

(3.8) Consider the bank database of Figure 3.19 where the primary keys are underlined. Construct the following SQL queries for the relational database.

(a) Find all customers of the bank who have an account but not a loan.

(b) Find the names of all customers who live on the same street and in the same city as "Smith".

(c) Find the names of all branches with customers who have an account in the bank and who live in 'Harrison'.

SOLN:

(a) Find

(select customer_name from depositor) except
(select customer_name from borrower)

(b) Select F.customer_name from

customer F join customer S using (customer_street,
customer_city) where S.customer_name = 'Smith'

(c) select distinct branch_name from
account natural join depositor natural join
customer where customer_city = 'Harrison'

(3.9) Consider the employee database of Figure
B.20, where the primary keys are underlined.
Give an expression SQL for each of the
following queries.

employee (employee_name, street, city)

works (employee_name, company_name, salary)

company (company_name, city)

manages (employee_name, manager_name)

(a) Find the names and cities of residence of
all employee who work for First Bank
corporation

(b) Find the names, street address, and cities of
residence of all employees who work for First
Bank corporation and earn more than \$10,000.

(c) Find all employees in the database who do not
work for First Bank Corporation.

- (d) Find all employees in the database who earn more than each employee of Small Bank Corporation.
- (e) Assume that, the companies may be located in several cities. Find all companies located in every city in which Small Bank Corporation.
- (f) Find the company that has most employees.
- (g) Find those companies which employees earn a higher salary, on average, than the other average salary at First Bank Corporation

Soln:

- (a) select e.employee-name, city from employee e, works w where w.company-name = 'First Bank Corporation' and w.employee-name = e.employee-name
- (b) select * from employee where employee-name in (select employee-name from works where company-name = 'First Bank Corporation' and salary > 1500).
- (c) select employee-name from works where company-name != 'First Bank Corporation'
- (d) select employee-name from works where salary > all (select salary from works where company-name = 'Small Bank Corporation')

(e) Select S.company-name from company S.
where not exists ((select city from company
where company-name = 'Small Bank Corporation')
except (select city from company T where
S.company-name = T.company-name)).

(f) Select company-name from works group by
company-name having count(distinct
employee-name) >= all (select count
distinct employee-name) from works
group by employee-name.

(g) Select company-name from works group by
company-name having avg(salary) > (select
avg(salary) from works where
company-name = 'First Bank Corporation')

(3.10) Consider the relational figure 3.20. Give
an expression in SQL for each of the
following queries.

(a) Modify the database so that Jones now
lives in Newton.

(b) Give all managers of First Bank Corporation
a 10 percent raise unless the salary becomes
greater than \$100,000; in such cases, give
only a 3 percent raise.

Sol:

- (a) update employee set city = 'Newton' where person-name = 'jones'
- (b) update works T set T.salary = T.salary *
(case
when (T.salary * 1.1 > 1000) then 1.03
else 1.1
)
- where T.employee-name in (select manager-name
from manages) and T.company-name = 'First Bank
Corporation'.

Exercises:

- (B.11) Write the following queries in sql, using the university schema.
- Find the name of all students who have taken at least one Comp. Sci. course; make sure there are no duplicate name in the result.
 - Find the ID's and name of all students who have not taken any course offering Spring 2009.
 - For each department, find the maximum salary of instructors in that department for at least one instructor.
 - Find the lowest, across all departments, of the per-department maximum salary computed by the preceding query.

SQl:

(a) select name from student natural join course
where course.dept = 'Comp. Sci'

(b) select id, name from student except
select id, name from student natural join
takes where year < 2009

(c) select dept, max(salary) from instructor
group by dept

(d) select min(maxsalary) from (select dept,
max(salary) as maxsalary from instructor
group by dept)

(3.12) Write the following queries in SQL, using
the university schema.

(a) Create a new course "CS-101", titled
"Weekly Seminar" with 0 credits.

(b) Create a section of this course in Autumn,
2009

(c) Enroll every student in the Comp. Sci.
department in the above section.

(d) Delete enrollments in the above section where
the student's name is Chavez.

(e) Delete the course CS-101. What will
happen if you run this delete statement
without first deleting offerings of this course?

(f) Delete all takes-tuples corresponding to any
section of any course with the word "database"
as a part of the title, ignore case when
matching the word with the title.

Soln:

- (a) insert into course values ('CS-001', 'Weekly Seminar', 'Comp. Sci.', 0)
- (b) insert into section values ('CS-001', 1, 'Autumn', 2009, null, null, null)
- (c) insert into takes
- select id, 'CS-001', 1, 'Autumn', 2009, null from student where dept_name = 'Comp. Sci'
- (d) delete from takes where course_id = 'CS-001' and section_id = 1 and year = 2009 and semester = 'Autumn' and id in (select id from student where name = 'Chavez')
- (e) delete from takes where course_id = 'CS-001'
delete from section where course_id = 'CS-001'
delete from course where course_id = 'CS-001'
- (f) delete from takes where course_id in
(select course_id from course where lower(title)
like '%database%')
- (3.13) Write SQL DDL corresponding to the schema in figure 3.18. Make any reasonable assumptions about data types and be sure to declare primary and foreign keys.

Solⁿ:

- (a) ~~create~~ create-table person (
- driver_id varchar(50),
 name varchar(50),
 address varchar(50),
 primary key (driver_id)).
- (b) create-table car (
- license varchar(50),
 model varchar(50),
 year integer,
 primary key (license))
- (c) create-table accident (
- report_number integer,
 date date,
 location varchar(50),
 primary key (report_number))
- (d) create-table owns (
- driver_id varchar(50),
 license varchar(50),
 primary key (driver_id, license),
 foreign key (driver_id) references person,
 foreign key (license) references car)
- (e) create-table participated (
- report_number integer,
 license varchar(50),
 driver_id varchar(50),
 damage_amount integer,

primary key (report-number, license).
 foreign key (license) references car,
 foreign key (report-number) references accident

(B.14) consider the insurance databases of Figure B.18, where the primary keys are underlined. construct the following SQL queries for this relational databases.

- (a) Find the number of accidents in which the cars belonging to 'John Smith' were involved.
- (b) Update the damage amount for the car with license number "AAB32000" in the accident with report number "AR2197" to \$3000.

SOL:

- (a) select count(*) from accident where exists (select * from participated owns person where owns.driver-id = person.driver-id, and person.name = 'John Smith' and own.license = participated.license and accident.report-number = participated.report-number)
- (b) update participated set damage-amount = 3000 where report-number = 'AR2197' and license = 'AAB32000')

(B.15) Consider the bank database of Figure B.19, where the primary keys are underlined. Construct the following SQL queries for this relational database.

branch (branch-name, branch-city, assets)

customer (customer-name, customer-street,
customer-city)

loan (loan-number, branch-name, amount)

borrower (customer-name, loan-number)

account (account-number, branch-name,
balance),

depositor (customer-name, account-number)

Fig: 3.19

(a)

(a) Find all customers who have an account at all the branches located in "Brooklyn".

(b) Find out the total sum of all loan amounts in the bank.

(c) Find the names of all branches that have assets greater than those of at least one branch located in "Brooklyn".

Soln:

(a) with branchcount as

(select count(*) branch where
branch-city = 'Brooklyn')

select customer-name from customer e

where branch_count = (select count(distinct
branch_name) from (customer natural join
depositor natural join account ~~as~~ natural join
branch) as d
where d.customer_name = e.customer_name)

(b) select sum(amount) from ban.

(c) select branch_name from branch
where assets > some(select assets from
branch where branch_city = 'Brooklyn')

(3.16) Consider the employee database of
figure 3-20, where the primary keys are
underlined.

employee (employee-name, street, city)

works (employee-name, company-name, salary)

company (company-name, city)

manages (employee-name, manager-name)

(a) Find the names of all ~~department~~ employees
who work for ~~first~~ Bank Corporation.

(b) Find all employees in the database who
live in the same cities as the company for
which they work.

(c) Find all employees in the database who live in
the same cities and on the same streets as do
their managers.

- (d) Find all employees who earn more than the average salary of all employees of their company.
- (e) Find the company that has the smallest payroll.

Soln:

- (a) select employee-name from works where company-name = 'First Bank Corporation'
- (b) select e.employee-name from employee e, works, company c
where e.employee-name = w.employee-n. and e.city = c.city and w.company-name = c.company-name
- (c) select P.employee-name from employee P, employee R, manager M
where P.employee-name = M.employee-name and M.manager-name = R.employee-name and P.street = R.street and P.city = R.city.
- (d) select employee-name from works +
where salary > (select avg(salary) from works s)

(e) select company-name from works group by company-name
having sum(salary) <= all (select sum(salary)
from works group by company-name)

(3.17) Consider the relation of figure 3.20.

- Give all employees of First Bank Corporation a 10 percent raise.
- Give all managers of First Bank Corporation a 10 percent raise.
- Delete all tuples in the works relation for employees of Small Bank Corporation.

SOP:

- update works set salary = salary * 1.1
where company-name = 'First Bank Corporation'
- update works set salary = salary * 1.1
where employee_name in (select manager_name
from manages) and company-name =
'First Bank Corporation'
- delete from works where company-name =
'Small Bank Corporation'.

(B.21) consider the library database of figure B.21.

members (memb_no, name, age)

book (isbn, title, author, publisher)

borrowed (memb_no, isbn, date)

Fig. B.21

- Print the names of members who have borrowed any book published by "Mc Graw-Hill".
- Print the names of members who have borrowed all books published by 'Mc Graw-Hill'
- For each publisher, print the names of members who have borrowed more than five books of that publisher.
- Print the average number of books borrowed per member. Take into account that if a member does not borrow any books, then that member does not appear in the borrowed relation at all.

Soln:

(a) Select name
from members m, book b, borrowed l,
where $m.memb_no = l.memb_no$
and $l.isbn = b.isbn$ and
 $b.publisher = 'Mc Graw Hill'$

(b) select distinct m.name
from member m where not exists
(select isbn from book where
publisher = 'McGrawHill')
except
(select isbn from borrowed l
where l.memb_no = m.memb_no))

(c) select publisher, name
from (select publisher, name, count(isbn)
from member m, book b, borrowed l
where m.memb_no = l.memb_no,
and l.isbn = b.isbn group by
publisher, name) as
membpub (publisher, name, count_books)
where count_books > 5.

(d) With mem count of
(select count(*) from member)
select count(*) / memcount from borrowed.

Chapter 04: Intermediate SQL

Practise Question:

(4.1) Write the following queries in SQL:

(a) Display a list of all instructors, showing their ID, name and the number of sections that they were taught. Your query should use outerjoin and should not use scalar queries.

(b) Write the same query above, but using a scalar subquery, without outerjoin.

(c) Display the list of all courses sections offered in Spring, 2010 along with the names of the instructors teaching the section. If a section has more than one instructors it should appear as many times in the result as it has instructors.

(d) Display the list of all departments, with the total number of instructors in each department, without using scalar subqueries.

Soln:

(a) Select ID, name, count(course_id, section_id, year semester) as number_of_sections from instructor natural left outer join teaches group by ID, name.

(b) Select ID, name (select count(*) as 'Number of sections' from teaches where T.id = I.id) from instructor I.

(c) select course_id, section_id, ID,
decode(name, NULL, '—', name)
from (section natural join left outer join teaches)
natural left outer join instructors
where semester = 'Spring' and year = 2010

(d) select dept_name, count(ID)
from department natural join left outer join instructors
group by dept_name

(4.2) Outer join expression can be computed in SQL
without using the SQL outer join.

- (a) select * from student natural left outer join takes
(b) select * from student natural full outer join takes.

Soln:

(a) select * from student natural join takes
union

select ID, name, dept_name, tot_cred, NULL, NULL, NULL
from student S1 where not exists

(select ID takes T1 where T1.id = S1.id)

(b) (select * from student natural join takes)
union

(select ID, name, dept_name, tot_cred, NULL,
NULL, NULL, NULL, NULL) from student S1

where not exists (select ID from takes
T1 where T1.id = S1.id)

union

(select ID, NULL, NULL, NULL, course_id, section_id,
semester, year, grade from takes T1
where not exists (select ID from student
where T1.id = s1.id))

(4.5) Show how to define the view student_grades
(ID, GPA) giving the grade-point average of each
student based on the query - recall that
we used a relation grade-points (grade_points)
the numeric points associated with a letter grade.

Soln: create view student_grades (ID, GPA) as

select ID, credit_points / decode (credit_sum, 0,
NULL, credit_sum)

from ((select ID, sum(decode (grade, NULL, 0, credits))
as credit_sum, sum (decode (grade, null, 0,
credits * points)) as credit_point

from (takes natural join course natural left
outerjoin grade_points group by ID)

Union

select ID, NULL as from student
where ID not in (select ID from takes))

(4.7) Consider the relational database of Figure 4.11.
Give an SQL DDL definition of this database.
Identify referential-integrity constraints that
should hold, and include them in the DDL Definition.

Soln:

Create table employee (

person-name char(20),

street char(30)

city char(30),

primary key (person-name))

Create table works

(person-name char(20),

company-name char(15),

salary integer,

primary key person-name,

foreign key person-name references employee)

Create table company

(company-name char(15),

city char(20),

primary key (company-name));

Create table manages

(person-name char(20),

manager-name char(20),

primary key (person-name),

foreign key (person-name) references employee);

(4.8) As discussed in section 4.4.7 complex check conditions and assertions subsection 4.4 we expect the constraints "an instructor cannot teach sections in two different classroom in a semester in the same time slot" to hold.

- (a) Write a SQL query that returns all combinations that violate the constraints.
(b) Write an SQL assertion enforce this constraint.

Soln:

- (a) select ID, name, section_id, semester, year, time-slot-id, count (distinct building, room-number) from instructor natural join teaches natural join section group by (ID, name, semester, year, time-slot-id) having count (building, room-number) > 1
- (b) create assertion check not exists
(select ID, name, section_id, semester, year, time-slot-id, count (distinct building, room-number) from instructor natural join teaches natural join section group by (ID, name, section_id, semester, year, time-slot-id) having count (building, room-number) > 1)

Exercises :

(4.12) For the database of figure 4.11, write a query to find those employees with no manager. Write your query using outer join with and then again using no outer join.

Soln:

a) select employee_name

from employee natural left outer join manages.

where manager_name is null

b) select employee_name from employee e
where not exists

(select employee_name from manages m

where e.employee_name = m.employee_name
and m.manager_name is not null)

(4.14) Show how to define a view in tot_credits
(year, num_credits), giving the total number of credits taken by students in each year.

Soln:

Create view tot_credits (year, tot_credits) as

(select year, sum(credits) ~~from~~

from takes natural join course

group by year)

(4.15) Show how to express the coalesce operation from exercise 4.10 using the case operation,

Soln:

Select

case Result

when (A_1 is not null) then A_1 ,

when (A_n is not null) then A_n ,

else null

end

from A

Chapter 06: Formal Relational Query Language

Practise exercises:

(6.1) Write the following queries in relational algebra, using the university schema.

- (a) Find the titles of courses in the Comp. Sci. department that have 3 credits.
- (b) Find the ID's of all students who were taught by an instructor named Ernestine.
- (c) Find the highest salary of any instructor.
- (d) Find all instructors earning the highest salary.
- (e) Find the enrollment of each section that was offered in Autumn, 2009.
- (f) Find the maximum enrollment, across all sections, in autumn 2009.
- (g) Find the sections that had the maximum enrollment in Autumn, 2009.

Soln:

- (a) $\Pi_{\text{title}} (\text{Dept-name} = \text{'comp.sci'} \wedge \text{credits} = 3 \text{ (course)})$
- (b) $\Pi_{\text{ID}} (\text{ID} = \text{'Einstei'} \text{ (takes} \bowtie \text{ P}_{t1} (\text{ID, course-id, section-id, semester, from}) \text{ teaches}))$
- (c) $\text{Gmax}(\text{salary})(\text{instructor})$
- (d) $\text{instructor} \bowtie (\text{Gmax}(\text{salary}) \text{ as salary} (\text{instructor}))$
- (e) $\text{course-id, section-id} \text{ Gcount}(x) \text{ as enrollment} (\text{year} = 2009 \wedge \text{semester} = \text{Autumn} (\text{takes}))$
- (f) $t_1 \leftarrow \text{course-id, section-id Gcount}(x) \text{ as enrollment}$
 $(\text{year} = 2009 \wedge \text{semester} = \text{Autumn} (\text{takes}))$
result = $\text{Gmax}(\text{enrollment}) \text{ as enrollment } (t_1)$
- (g) $t_2 \leftarrow \text{Gmax}(\text{enrollment}) \text{ as enrollment } (t_1)$
result = $t_1 \bowtie t_2$

(6.2) Consider the relational database of Figure 6.22, where the primary keys are underlined. Give an expression in the relational algebra to express each of the following queries:

- (a) Find the names of all employees who live in the same city and on the same street as do their managers.
- (b) Find the names of all employees in this database who do not work for "First Bank Corporation".

(c) Find the names of all employees who earn more than every employee of "Small Bank Corporation".

Soln:

(a) $\exists \text{person_name} (\text{Employee} \bowtie \text{manages})$

$\forall (\text{manager_name} = \text{employee2.person_name} \wedge \text{employee.street} = \text{employee2.street}) (\text{Employee}_2(\text{employee}))$

(b) $\exists \text{person_name} (\text{Company_name} = \text{First Bank Corporation} (\text{works}))$

(c) $\exists \text{person_name} (\text{works}) - \exists \text{works}.\text{person_name} (\text{works}$

$\forall \text{works}.\text{salary} \leq \text{works2.salary} \wedge \text{works2.company_name} = \text{"Small Bank Corporation"} \text{ } \exists \text{works2}(\text{works}))$

(6.3) The natural outer join operations extend the natural join operation so that tuples from the participating relations are not lost in the result of the join.

Describe how the theta-join operation can be extended so that tuples from the left, right, or both relations are not lost from the result of a theta join.

Soln:

(a) The left outer theta join of $\pi(R)$ and $s(S)$

$(r \bowtie_{\theta} s)$ can be defined as $(r \bowtie_{\theta} s) \cup$

$(r - \pi_r(r \bowtie_{\theta} s)) \times (\text{null}, \text{null}, \dots, \text{null})$.

The tuple of nulls is of size equal to the number of attributes in s .

(b) The right outer theta join of $r(R)$ and $s(S)$
 $(r \bowtie_{\theta} s)$ can be defined as $(r \bowtie_{\theta} s) \cup ((\text{null} \dots \text{null}) \times (s - \text{TT}_s(r \bowtie_{\theta} s)))$

(6.7) Let $R = (A, B)$ and $S = (A, C)$ and $r(R)$ and $s(S)$ be relations. Write expressions in relational algebra for each of the following queries

(a) $\{ \langle a \rangle \mid \exists b (\langle a, b \rangle \in r \wedge b = 7) \}$

(b) $\{ \langle a, b, c \rangle \mid \langle a, b \rangle \in r \wedge \langle a, c \rangle \in s \}$

(c) $\{ \langle a \rangle \mid \exists c (\langle a, c \rangle \in s \wedge \exists b_1, b_2 (\langle a, b_1 \rangle \in r \wedge \langle a, b_2 \rangle \in r \wedge b_1 > b_2)) \}$

Soln:

(a) $\Pi_A (G_{B=7}(r))$

(b) $r \bowtie s$

(c) $\Pi_A (S \bowtie (\Pi_{R,A} (G_{R,B} \rightarrow d.b(r \times \rho_B(r)))))$

(6.8) Consider the relational database of figure 6.22 where the primary keys are underlined. derive an expression in tuple relational calculus for each of the following question.

(a) Find all employees who work directly for "Jones".

(b) Find all cities residence of all employees who work directly for 'Jones'.

(c) Find the name of the manager of the manager of 'Jones'.

Soln:

(a) $\{ \text{ } \} \uplus \exists m \in \text{manages} \ (\text{ } \llcorner \text{ [person_name]} = m \text{ [person_name]})$
 $\wedge m \text{ [manager_name]} = \text{'Jones'} \}$

(b) $\{ \text{ } \} \uplus \exists m \in \text{manages} \ \exists s \in \text{employee} ($
 $e \text{ [person_name]} = m \text{ [person_name]} \wedge$
 $m \text{ [manager_name]} = \text{'Jones'} \wedge$
 $t \text{ [city]} = e \text{ [city]}) \}$

(c) $\{ \text{ } \} \uplus \exists m_1 \in \text{manages} \ \exists m_2 \in \text{manages} ($
 $m_1 \text{ [manager_name]} = m_2 \text{ [person_name]}$
 $\wedge m_1 \text{ [person_name]} = \text{'Jones'}$
 $\wedge t \text{ [manager_name]} = m_2 \text{ [manager_name]}) \}$

Exercise:

(6.10) Write the following queries in relational algebra, using the university schema.

- Find the names of all students who have taken at least one Comp. Sci course.
- Find the IDs and names of all students who have not taken any course Spring, 2009.
- For each department, find the maximum salary of instructor in the department.

(d) Find the lowest, across all departments, of the per-department maximum salary computed by the preceding query.

Soln:

(a) $\Pi_{\text{name}} (\text{student} \bowtie \text{takes} \bowtie \Pi_{\text{course-id}} (\text{dept-name} = \text{'comp.sci'}) \text{course}))$

(b) $\Pi_{\text{dept-name}} (\text{Student}) = \Pi_{\text{id-name}} (\text{Byear} < 2009) (\text{student} \bowtie \text{takes})$

(c) $\text{dept-name} \sqcup_{\max(\text{salary})} (\text{instructor})$

(d) $\sqcup_{\min(\max(\text{salary}))} (\text{dept-name} \sqcup_{\max(\text{salary})} \text{as maxsal} (\text{instructor}))$

(6.11) Consider the relational database 6.22, where the primary keys are underlined. Give an expression in the relational database algebra to express each of the following queries.

(a) employee (person-name, street, city)

works (person-name, company-name, salary)

company (Company-name, city)

manages (person-name, manager-name)

Fig: 6.22

(a) Find the names of all employees who work for "First Bank Corporation".

(b) Find the names and cities of residence of all employees who work for "First Bank Corporation".

- (b) Find the names and cities of residence of all employees who work for "First Bank Corporation".
- (c) Find the names, street address and cities of residence of all employees who work for "First Bank Corporation" and earn more than \$10,000.
- (d) Find the names of all employees in this database who live in the same city as the company for which they work.
- (e) Assume the companies may be located in several cities. Find all companies located in every city in which "Small Bank Corporation" is located.

Solⁿ:

- (a) $\Pi_{\text{person_name}} (\text{Company-name} = \text{'First Bank Corporation'})^{(\text{works})}$
- (b) $\Pi_{\text{person_name}, \text{city}} (\text{Employee} \bowtie (\text{Company-name} = \text{'First Bank Corporation'})^{(\text{works})})$
- (c) $\Pi_{\text{person_name}} (\text{Employee} \bowtie \text{works} \bowtie \text{Company})$
- (c) $\Pi_{\text{person_name}, \text{street}, \text{city}} ((\text{Company-name} = \text{'First Bank Corporation'}) \wedge \text{Salary} > 10000) \text{ works} \bowtie \text{Employee}$

(e) $\Pi_{\text{company_name}} (\text{company} \div$
 $(\Pi_{\text{city}} (\text{Company_Name} = \text{'Small Bank Corporation'}) \text{company}))$

(6.12) Using the university example, write relational-algebra queries to find the course sections taught by more than one instructor in the following ways:

(a) Using an aggregate function.

(b) Without using any aggregate function.

Soln:

(a) $\exists \text{instruct} > 1 (\text{course_id}, \text{section_id},$
 $\text{year}, \text{semester} \text{ } \text{Gcount}(\text{*}) \text{ as instruct } \text{ (Teaches)})$

(b) $\Pi_{\text{course_id}, \text{section_id}, \text{year}, \text{semester}} (\text{GID} <$
 $\text{ID}_2 \text{ (Takes} \bowtie \text{ Takes1 (ID}_2, \text{course_id}, \text{section_id}, \text{year},$
 $\text{semester}) \text{ (Takes)))}$

(6.13) Consider the Relational database of Figure 6.22. Give a relational-algebra expression for each of the following queries:

- (a) Find the company with the most employees.
(b) Find the company with smallest payroll.

(c) Find those companies whose employees earn a higher salary on average at First Bank Corporation.

Solⁿ:

(a) $t_1 \leftarrow \text{company-name} \setminus \text{count-distinct}(\text{person-name})$ (works)
 $t_2 \leftarrow \text{g}_{\max}(\text{max-employees}) (\text{Pcompany-strength}(\text{company-name}, \text{num-employees})(t_1))$

$\Pi_{\text{company-name}}(\text{P}_{t_3}(\text{company-name}, \text{num-employees})(t_1) \bowtie \text{P}_{t_4}(\text{num-employees})(t_2))$

(b) $t_1 \leftarrow \text{company-name} \setminus \text{sum}(\text{salary})$ (works)
 $t_2 \leftarrow \text{g}_{\min}(\text{payroll}) (\text{Pcompany-payroll}(\text{company-name}, \text{payroll}))$

$\Pi_{\text{company-name}}(\text{P}_{t_3}(\text{company-name}, \text{payroll})(t_1) \bowtie \text{P}_{t_4}(\text{payroll})(t_2))$

(c) $t_1 \leftarrow \text{company-name} \setminus \text{avg}(\text{salary})$ (works)

$t_2 \leftarrow \text{C}_{\text{company-name}} = \text{First Bank Corporation}$ (t1)

$\Pi_{t_3. \text{company-name}}(\text{P}_{t_3}(\text{company-name}, \text{avg-salary})(t_1))$

$\bowtie t_3. \text{avg-salary} > \text{first_bank.avg-salary} (\text{P}_{\text{first_bank}}(\text{company-name}, \text{avg-salary})(t_2))$

(6.14) Consider the following relational schema for a library:

member (memb_no, name, dob)
books (isbn, title, authors, publisher)
borrowed (memb_no, isbn, date)

B6.

- (a) Find the names of members who have borrowed any book published by 'McGraw-Hills'
- (b) Find the name of members who have borrowed all books published by 'McGraw-Hills'.
- (c) Find the name and membership number of members who have borrowed more than five different books published by McGraw-Hills'
- (d) For each publisher, find the name of books borrowed per member(s) who have borrowed more than five books of that publisher.

SOL:

- (a) $t_1 \leftarrow \text{isbn} (\sigma_{\text{publisher} = \text{'McGraw-Hills'}}(\text{books}))$
 $\text{Tname} ((\text{member} \bowtie \text{borrowed}) \bowtie t_1)$
- (b) $t_1 \leftarrow \text{isbn} (\sigma_{\text{publisher} = \text{'McGraw-Hills'}}(\text{books}))$
 $\text{Tname, isbn } (\text{member} \bowtie \text{borrowed}) \div t_1$
- (c) $t_1 \leftarrow \text{member} \bowtie \text{borrowed} \bowtie (\sigma_{\text{publisher} = \text{'McGraw-Hill'}}(\text{books}))$
 $\text{Tname} (\sigma_{\text{countisbn} > 5} (\text{count-distinct(isbn)} \text{ or } \text{countisbn}(t_1)))$
- (d) $t_1 \leftarrow \text{member} \bowtie \text{borrowed} \bowtie \text{books}$
 $\text{Tpublisher, name} (\sigma_{\text{countisbn} > 5} (\text{publisher, memb-no, count-distinct of countisbn}(t_1)))$
- (e.16) Let $R = (A, B)$ and $S = (A, C)$ and let $\pi(R)$ and $s(S)$ be rotations. Write relational-algebraic expressions equivalent to the following domain-relational-calculus:

- (a) $\{a \mid \exists b (a, b \in r \wedge b = 17)\}$
- (b) $\{a-b, c \mid a-b \in r \wedge (a, c) \in s\}$
- (c) $\{a \mid \exists b (a, b \in r) \vee \forall c (\exists d (d, c \in s)$
 $\Rightarrow a, c \in s)\}$
- (d) $\{a \mid \exists c (a, c \in s \wedge \exists b_1, b_2 (a, b_1 \in r \wedge$
 $a, b_2 \in r \wedge b_1 > b_2))\}$

Soln:

- (a) $\text{TT}_A (\sigma_B = 17(r))$
- (b) $r \bowtie s$
- (c) $\text{TT}_A(r) \cup (r \div \sigma_B(\text{TC}(s)))$
- (d) $\text{TT}_{T_A} ((r \bowtie s) \wedge_c = r.A \wedge r.B \wedge r_2.B \quad (\text{Pr}_2(r)))$

(6.17) Repeat exercise 6.16, writing SQL queries instead of relational-algebra expressions.

Soln:

- (a) select a from r where b = 17
- (b) select a, b, c from r, s where r.a = s.a
- (c) $\text{TT}(\text{select a from r}) \text{ union } (\text{select a from s})$
- (d) select a from r as r1, r as r2, s
 where r1.a = s.a and r2.a = s.c and r1.b > r2.b

(6.18) Let $R = (A, B)$ and $S = (A, C)$ and let $r(R)$ and $s(S)$ be relations. Using the special constants null, write tuple-relational-calculus expressions equivalent to each of the following:

$$(a) r \bowtie s$$

$$(b) r \bowtie s$$

$$(c) r \bowtie s$$

Soln:

$$(a) \{d \mid \exists r \in R \exists s \in S (r[A] = s[A] \wedge t[A] = r[B] \wedge t[B] = r[B] \wedge t[C] = s[C]) \vee \exists s \in S (\neg \exists r \in R (r[A] = s[A]) \wedge t[A] = s[A] \wedge t[C] = s[C] \wedge t[B] = \text{null})\}$$

$$(b) \{d \mid \exists r \in R \exists s \in S (r[A] = s[A] \wedge t[A] = r[A] \wedge t[B] = r[B] \wedge t[C] = s[C]) \vee \exists r \in R (\neg \exists s \in S (r[A] = s[A]) \wedge t[A] = r[A] \wedge t[B] = r[B] \wedge t[C] = \text{null}) \vee \exists s \in S (\neg \exists r \in R (r[A] = s[A]) \wedge t[A] = s[A] \wedge t[C] = s[C] \wedge t[B] = \text{null})\}$$

$$(c) \{d \mid \exists r \in R \exists s \in S (r[A] = s[A] \wedge t[A] = r[A] \wedge t[B] = r[B] \wedge t[C] = s[C]) \vee \exists r \in R (\neg \exists s \in S (r[A] = s[A]) \wedge t[A] = r[A] \wedge t[B] = r[B] \wedge t[C] = \text{null})\}$$

(6.19) Give a tuple-relational-calculus expressions to find the maximum value in T relational $\pi(A)$.

Soln: $\{ \langle a \rangle \mid \langle a \rangle \in \pi \wedge \langle b \rangle \in R \text{ and } a > b \}$