

Some repeated questions

👤 Created by	🅑 Borhan
🕒 Last edited time	@December 9, 2024 9:48 PM
⋮ Tag	

A transaction can be defined as a group of tasks. A transaction is a single collection of operations that forms a single logical unit of work.

- **Atomicity** – This property states that a transaction must be treated as an **atomic unit**, that is, either all of its operations are **executed or none**. There must be no state in a database where a transaction is left partially completed. States should be defined either before the execution of the transaction or after the execution/abortion/failure of the transaction.
- **Consistency** – The database must **remain in a consistent state after any transaction**. No transaction should have any **adverse effect** on the data residing in the database. If the database was in a consistent state before the execution of a transaction, it must remain consistent after the execution of the transaction as well.
- **Durability** – The database should be **durable enough to hold all its latest updates** even if the **system fails** or restarts. If a transaction updates a chunk of data in a database and commits, then the database will **hold the modified data**. If a transaction commits but the **system fails** before the data could be written on to the disk, then that data **will be updated once the system springs** back into action.
- **Isolation** – In a database system where more than one transaction are being **executed simultaneously** and in **parallel**, the property of isolation states that all the transactions will be carried out and executed as if it is the **only transaction in the system**. **No transaction will affect the existence of any other transaction**.

Serializability

Schedule – A chronological **execution sequence of a transaction** is called a schedule. A schedule can have **many transactions** in it, each comprising of a

number of instructions/tasks.

A concurrent schedule is serializable if its outcome is the same as **some serial schedule of the same transactions**.

Key Concepts of Serializability

1. **Serial Schedule:** A schedule where **transactions are executed sequentially**, one after another, without interleaving. For example:
 - Transaction T1 executes entirely before Transaction T2 begins.
2. **Concurrent Schedule:** A schedule where the **operations of transactions are interleaved**. This can **improve performance** but may lead to **conflicts** or inconsistencies.

Result Equivalence

If two schedules produce the same result after execution, they are said to be result equivalent.

Conflict Equivalence


Two schedules would be **conflicting** if they have the following properties –

- Both belong to separate transactions.
- Both accesses the same data item.
- At least one of them is "write" operation.

Two schedules having multiple transactions with conflicting operations are said to be **conflict equivalent if and only if** –

- Both the schedules contain the same set of Transactions.
- The order of conflicting pairs of operation is maintained in both the schedules.

A schedule is called conflict serializability if after **swapping of non-conflicting operations**, it can transform into a **serial schedule**.



- **Active** – In this state, **the transaction is being executed**. This is the **initial state** of every transaction.
- **Partially Committed** – When a transaction **executes its final operation**, it is said to be in a partially committed state.
- **Failed** – A transaction is said to be in a failed state if any of the checks made by the database recovery **system fails**. A failed transaction **can no longer proceed** further.
- **Aborted** – If any of the checks fails and the transaction has reached a failed state, then the **recovery manager rolls back all its write operations** on the database to bring the database back to its original state where **it was prior to the execution of the transaction**. Transactions in this state are called aborted. The database recovery module can select one of the two operations after a transaction aborts –
 - Re-start the transaction
 - Kill the transaction
- **Committed** – If a transaction **executes all its operations successfully**, it is said to be committed. All its effects are now **permanently established** on the database system.

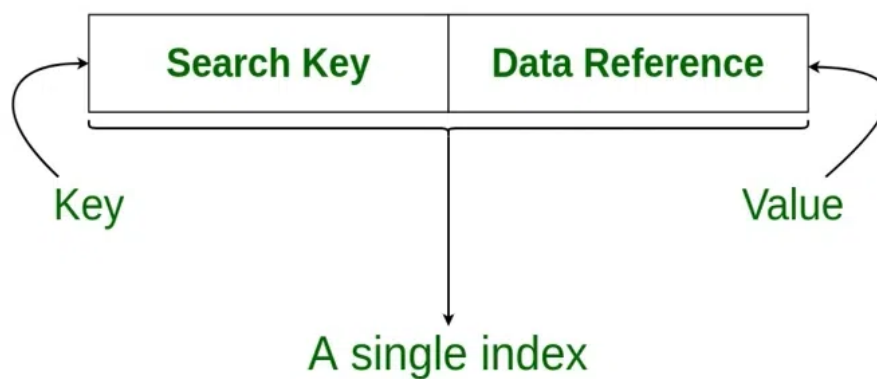
Indexing

Indexing is technique for **improving database performance** by **reducing the number** of disk accesses necessary when a **query** is run.

Why Indexing is needed-

1. Speed up access
2. Minimizing number of disk access
3. Less I/O Operation

Structure of an Index in Database



EG

Deadlock

The Deadlock is a condition in a **multi-user database** environment where transactions are unable to complete because they are **each waiting for the resources held by other transactions**.